

# EuroSAT Classification with Swin Transformer and Advanced Techniques (Version 1)

Documented Code

February 14, 2025

Section 1 of this document provides an annotated listing of the code used for optimized Swin Transformer for EuroSAT Classification with advanced techniques such as Mixup, CosineAnnealingWarmRestarts, GradScaler for mixed precision, and a custom parameter scheduling function. The code is designed to run within a Python environment (e.g., Kaggle notebooks, Google Colab) that supports PyTorch, timm, and scikit-learn.

Section 2 provides the code to train each of the 5 ViT architecture models, namely DeiT, YOLO, Swin, PVT and MAE. This was used to select the baseline model used in Section 1 above.

## 1 Optimized Swin Transformer

Below is a comprehensive code listing for training a *Swin Transformer* on the EuroSAT dataset using advanced techniques:

- **Data Augmentation** via **RandAugment**, random flips, random erasing, etc.
- **Mixup** to enhance robustness.
- **Layer-wise LR** scheduling for finer control over learning rates in deeper blocks.
- **CosineAnnealingWarmRestarts** for cyclical LR patterns.
- **SWA (Stochastic Weight Averaging)** to stabilize and improve final performance.
- **Fine-Tuning** stage that freezes earlier layers and adjusts LR for the final classification head.

---

Listing 1: EuroSAT + Swin Base Patch4 Window7 224 with RandAugment, Mixup, SWA, and Fine-Tuning

```
1 import os
2 import glob
3 import pandas as pd
4 import numpy as np
5
6 from sklearn.model_selection import train_test_split
```

```

7 from sklearn.preprocessing import LabelEncoder
8 from sklearn.metrics import classification_report, confusion_matrix
9
10 import seaborn as sns
11 import matplotlib.pyplot as plt
12
13 import torch
14 import torch.nn as nn
15 import torch.optim as optim
16 from torch.utils.data import Dataset, DataLoader
17 from torchvision import transforms
18 from PIL import Image
19
20 # Additional transforms
21 from torchvision.transforms import RandAugment, RandomErasing
22
23 # timm (SOTA models) + Mixup
24 from timm import create_model
25 from timm.data.mixup import Mixup
26
27 # PyTorch optimization
28 from torch.optim import AdamW
29 from torch.optim.lr_scheduler import CosineAnnealingWarmRestarts
30
31 # AMP (mixed precision) and GradScaler
32 from torch.cuda.amp import autocast, GradScaler
33 # SWA from PyTorch
34 from torch.optim.swa_utils import AveragedModel, SWALR
35
36 # Seeds for reproducibility
37 torch.manual_seed(42)
38 np.random.seed(42)
39
40 #####
41 # 1. DATASET PREPARATION
42 #####
43 data_dir = '/kaggle/input/eurosat10-classes/EuroSAT_RGB/'
44 # 'data_dir' has subfolders named after land-cover classes (e.g., "AnnualCrop", "Forest",
45 etc.)
46
47 # Gather all .jpg paths
48 image_paths = glob.glob(os.path.join(data_dir, '*', '*.jpg'))
49 labels = [os.path.basename(os.path.dirname(path)) for path in image_paths]
50
51 # Make a DataFrame
52 df = pd.DataFrame({'image_path': image_paths, 'label': labels})
53
54 # Encode label strings into integers
55 le = LabelEncoder()
56 df['label_enc'] = le.fit_transform(df['label'])
57
58 # Stratified split: 80% train, 20% validation
59 train_df, val_df = train_test_split(
    df,

```

```

60     test_size=0.2,
61     stratify=df['label_enc'],
62     random_state=42
63 )
64
65 print(f"Total images: {len(df)}")
66 print(f"Training images: {len(train_df)}")
67 print(f"Validation images: {len(val_df)}")
68
69 #####
70 # 2. IMAGE TRANSFORMS
71 #####
72 # Mean/Std for EuroSAT (approx)
73 mean = [0.3444, 0.3809, 0.4082]
74 std = [0.1829, 0.1603, 0.1321]
75
76 train_transforms = transforms.Compose([
77     transforms.Resize(256),
78     transforms.RandomResizedCrop(224, scale=(0.8, 1.0)),
79     transforms.RandomHorizontalFlip(),
80     transforms.RandomVerticalFlip(),
81     RandAugment(num_ops=2, magnitude=7), # apply RandAugment
82     transforms.ToTensor(),
83     transforms.Normalize(mean=mean, std=std),
84     RandomErasing(p=0.5, scale=(0.02, 0.2)) # random erasing
85 ])
86
87 val_transforms = transforms.Compose([
88     transforms.Resize(224),
89     transforms.CenterCrop(224),
90     transforms.ToTensor(),
91     transforms.Normalize(mean=mean, std=std),
92 ])
93
94 #####
95 # 3. CUSTOM DATASET
96 #####
97 class EuroSATDataset(Dataset):
98     """
99     A simple Dataset class that:
100     - Uses a DataFrame with 'image_path' and 'label_enc'
101     - Loads images from disk, applies transforms
102     - Returns (image_tensor, label_index)
103     """
104     def __init__(self, df, transforms):
105         self.df = df.reset_index(drop=True)
106         self.transforms = transforms
107
108     def __len__(self):
109         return len(self.df)
110
111     def __getitem__(self, idx):
112         image_path = self.df.loc[idx, 'image_path']
113         label_enc = self.df.loc[idx, 'label_enc']

```

```

114         image = Image.open(image_path).convert('RGB')
115         image = self.transforms(image)
116         return image, label_enc
117
118     # Build Dataset + Dataloader
119     train_dataset = EuroSATDataset(train_df, train_transforms)
120     val_dataset = EuroSATDataset(val_df, val_transforms)
121
122     batch_size = 64
123     train_loader = DataLoader(
124         train_dataset,
125         batch_size=batch_size,
126         shuffle=True,
127         num_workers=4,
128         pin_memory=True
129     )
130     val_loader = DataLoader(
131         val_dataset,
132         batch_size=batch_size,
133         shuffle=False,
134         num_workers=4,
135         pin_memory=True
136     )
137
138     #####
139     # 4. CREATE SWIN MODEL
140     #####
141     model = create_model(
142         'swin_base_patch4_window7_224',
143         pretrained=True,
144         num_classes=10,
145         drop_rate=0.0,
146         drop_path_rate=0.1
147     )
148     # Cross-entropy with label smoothing
149     criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
150
151     #####
152     # 5. PARAM GROUPS + COSINE RESTARTS
153     #####
154     def get_param_groups(model, base_lr, weight_decay):
155         """
156         Build parameter groups for layer-wise LR scheduling.
157         We'll downscale LR for deeper blocks by 0.95^(layer_index).
158         """
159         no_weight_decay = model.no_weight_decay()
160         param_groups = {}
161
162         for name, param in model.named_parameters():
163             if not param.requires_grad:
164                 continue
165             group_name = 'layer_0'
166             if 'blocks' in name:
167                 block_num = int(name.split('.')[1])

```

```

168         group_name = f'layer_{block_num + 1}'
169     elif 'cls_token' in name or 'pos_embed' in name:
170         group_name = 'layer_0'
171     else:
172         group_name = 'layer_0'
173
174     if group_name not in param_groups:
175         param_groups[group_name] = {
176             'params': [],
177             'weight_decay': weight_decay,
178             'lr': base_lr
179         }
180     param_groups[group_name]['params'].append(param)
181
182     param_groups_list = []
183     num_layers = len(param_groups)
184     for i, (group_name, group) in enumerate(
185         sorted(param_groups.items(), key=lambda x: x[0])
186     ):
187         group['lr'] = base_lr * (0.95 ** (num_layers - i - 1))
188         param_groups_list.append(group)
189
190     return param_groups_list
191
192 base_lr = 3e-5
193 weight_decay = 0.01
194 optimizer = AdamW(get_param_groups(model, base_lr, weight_decay))
195
196 scheduler = CosineAnnealingWarmRestarts(optimizer, T_0=10, T_mult=1)
197
198 #####
199 # 6. MIXUP
200 #####
201 mixup_fn = Mixup(
202     mixup_alpha=0.8,
203     cutmix_alpha=1.0,
204     prob=1.0,
205     switch_prob=0.5,
206     mode='batch',
207     label_smoothing=0.1,
208     num_classes=10
209 )
210
211 #####
212 # 7. AMP / GRAD SCALER
213 #####
214 scaler = GradScaler()
215
216 # Move model to GPU
217 model = model.cuda()
218
219 #####
220 # 8. SWA SETUP
221 #####

```

```

222 swa_model = AveragedModel(model)
223 swa_start_epoch = 25
224 swa_scheduler = SWALR(
225     optimizer,
226     swa_lr=base_lr,
227     anneal_epochs=5,
228     anneal_strategy="cos"
229 )
230
231 #####
232 # 9. TRAINING (MIXUP, AMP, SWA)
233 #####
234 def train_one_epoch(epoch, model, dataloader, optimizer, criterion, scheduler, mixup_fn):
235     model.train()
236     running_loss = 0.0
237     total_samples = 0
238
239     for images, labels in dataloader:
240         images = images.cuda(non_blocking=True)
241         labels = labels.cuda(non_blocking=True)
242
243         # Apply mixup
244         images, labels = mixup_fn(images, labels)
245
246         optimizer.zero_grad()
247         with autocast(device_type='cuda', dtype=torch.float16):
248             outputs = model(images)
249             loss = criterion(outputs, labels)
250
251         scaler.scale(loss).backward()
252         nn.utils.clip_grad_norm_(model.parameters(), max_norm=5.0)
253         scaler.step(optimizer)
254         scaler.update()
255
256         running_loss += loss.item() * images.size(0)
257         total_samples += images.size(0)
258
259     epoch_loss = running_loss / total_samples
260
261     if epoch < swa_start_epoch:
262         scheduler.step()
263
264     print(f"[Epoch {epoch} | Train] Loss: {epoch_loss:.4f}")
265     return epoch_loss
266
267 def validate(model, dataloader, criterion):
268     model.eval()
269     running_loss = 0.0
270     correct = 0
271     total = 0
272     all_preds = []
273     all_labels = []
274
275     with torch.no_grad():

```

```

276     for images, labels in dataloader:
277         images = images.cuda(non_blocking=True)
278         labels = labels.cuda(non_blocking=True)
279
280         with autocast(device_type='cuda', dtype=torch.float16):
281             outputs = model(images)
282             loss = criterion(outputs, labels)
283
284             running_loss += loss.item() * images.size(0)
285             _, predicted = outputs.max(1)
286             total += labels.size(0)
287             correct += predicted.eq(labels).sum().item()
288
289             all_preds.extend(predicted.cpu().numpy())
290             all_labels.extend(labels.cpu().numpy())
291
292     epoch_loss = running_loss / total
293     epoch_acc = correct / total
294
295     print("Classification Report:")
296     print(classification_report(all_labels, all_preds, target_names=le.classes_, digits
297                                =4))
298
299     cm = confusion_matrix(all_labels, all_preds)
300     print(f"[Validation] Loss: {epoch_loss:.4f}, Accuracy: {epoch_acc:.4f}")
301     return epoch_loss, epoch_acc
302
303 #####
304 # 10. MAIN TRAIN LOOP
305 #####
306 num_epochs = 30
307 best_acc = 0.0
308
309 for epoch in range(1, num_epochs + 1):
310     print(f"Epoch {epoch}/{num_epochs}")
311
312     train_loss = train_one_epoch(
313         epoch, model, train_loader, optimizer, criterion, scheduler, mixup_fn
314     )
315
316     # SWA updates after 'swa_start_epoch'
317     if epoch >= swa_start_epoch:
318         swa_model.update_parameters(model)
319         swa_scheduler.step()
320
321     val_loss, val_acc = validate(model, val_loader, criterion)
322
323     if val_acc > best_acc:
324         best_acc = val_acc
325         torch.save(model.state_dict(), 'best_vit_model.pth')
326         print("Saved Best Model!")
327
328     print('-' * 40)

```

```

329 # Update BN for SWA model
330 torch.optim.swa_utils.update_bn(train_loader, swa_model, device='cuda')
331
332 # Save SWA model
333 torch.save(swa_model.state_dict(), 'swa_vit_model.pth')
334 print("SWA Model Saved!")
335
336 #####
337 # 11. FINE-TUNING PHASE
338 #####
339 import torch
340 import torch.nn as nn
341 from torch.cuda.amp import autocast, GradScaler
342 from torch.optim import AdamW
343 from torch.optim.lr_scheduler import CosineAnnealingWarmRestarts
344 from timm import create_model
345 from sklearn.metrics import classification_report, confusion_matrix
346 import matplotlib.pyplot as plt
347 import seaborn as sns
348
349 # Re-create the same model structure, but we'll load our best or SWA weights
350 model_finetune = create_model(
351     'swin_base_patch4_window7_224',
352     pretrained=False,
353     num_classes=10
354 )
355 model_finetune.load_state_dict(torch.load('best_vit_model.pth'))
356 model_finetune = model_finetune.cuda()
357
358 # Optionally freeze layers except final head
359 for name, param in model_finetune.named_parameters():
360     if 'head' not in name:
361         param.requires_grad = False
362
363 finetune_lr = 1e-5
364 finetune_weight_decay = 1e-4
365 finetune_optimizer = AdamW(
366     filter(lambda p: p.requires_grad, model_finetune.parameters()),
367     lr=finetune_lr,
368     weight_decay=finetune_weight_decay
369 )
370 finetune_scheduler = CosineAnnealingWarmRestarts(finetune_optimizer, T_0=5, T_mult=1)
371
372 finetune_criterion = nn.CrossEntropyLoss(label_smoothing=0.1)
373 finetune_scaler = GradScaler()
374
375 def finetune_one_epoch(epoch, model, dataloader, optimizer, criterion, scheduler):
376     model.train()
377     running_loss = 0.0
378     total_samples = 0
379
380     for images, labels in dataloader:
381         images = images.cuda(non_blocking=True)
382         labels = labels.cuda(non_blocking=True)

```



```

383
384     optimizer.zero_grad()
385     with autocast(device_type='cuda', dtype=torch.float16):
386         outputs = model(images)
387         loss = criterion(outputs, labels)
388
389     finetune_scaler.scale(loss).backward()
390     nn.utils.clip_grad_norm_(model.parameters(), 5.0)
391     finetune_scaler.step(optimizer)
392     finetune_scaler.update()
393
394     running_loss += loss.item() * images.size(0)
395     total_samples += images.size(0)
396
397     epoch_loss = running_loss / total_samples
398     scheduler.step()
399     print(f"[Fine-Tune Epoch {epoch}] Loss: {epoch_loss:.4f}")
400     return epoch_loss
401
402 def validate_finetune(model, dataloader, criterion):
403     model.eval()
404     running_loss = 0.0
405     correct = 0
406     total = 0
407     all_preds = []
408     all_labels = []
409
410     with torch.no_grad():
411         for images, labels in dataloader:
412             images = images.cuda(non_blocking=True)
413             labels = labels.cuda(non_blocking=True)
414
415             with autocast(device_type='cuda', dtype=torch.float16):
416                 outputs = model(images)
417                 loss = criterion(outputs, labels)
418
419             running_loss += loss.item() * images.size(0)
420             _, predicted = outputs.max(1)
421             total += labels.size(0)
422             correct += predicted.eq(labels).sum().item()
423
424             all_preds.extend(predicted.cpu().numpy())
425             all_labels.extend(labels.cpu().numpy())
426
427     epoch_loss = running_loss / total
428     epoch_acc = correct / total
429     print(f"[Fine-Tune Validation] Loss: {epoch_loss:.4f}, Accuracy: {epoch_acc:.4f}")
430
431     final_report = classification_report(all_labels, all_preds, target_names=le.classes_,
432                                         digits=4)
433     print("\nFinal Classification Report (Fine-Tuned Model):")
434     print(final_report)
435
436     cm = confusion_matrix(all_labels, all_preds)

```

```

436 plt.figure(figsize=(8,6))
437 sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
438             xticklabels=le.classes_, yticklabels=le.classes_)
439 plt.title("Confusion Matrix (Fine-Tuned Model)")
440 plt.ylabel('Actual')
441 plt.xlabel('Predicted')
442 plt.show()
443
444 return epoch_loss, epoch_acc
445
446 finetune_epochs = 10
447 best_acc_ft = 0.0
448
449 for ep in range(1, finetune_epochs+1):
450     train_loss_ft = finetune_one_epoch(
451         ep, model_finetune, train_loader, finetune_optimizer, finetune_criterion,
452         finetune_scheduler
453     )
454     val_loss_ft, val_acc_ft = validate_finetune(model_finetune, val_loader,
455         finetune_criterion)
456
457     if val_acc_ft > best_acc_ft:
458         best_acc_ft = val_acc_ft
459         torch.save(model_finetune.state_dict(), "best_vit_model_finetuned.pth")
460         print("Saved Best Fine-Tuned Model!\n")

```

## 2 Training 5 different ViT models

Listing 2: Downloading and Preparing EuroSAT Data

```

1 !wget http://madm.dfki.de/files/sentinel/EuroSAT.zip
2 !unzip EuroSAT.zip -d EuroSAT/
3
4 # Explanation:
5 # 1) We retrieve the EuroSAT dataset from dfki.de.
6 # 2) We unzip the dataset contents into a directory named 'EuroSAT/'.
7 # This yields subfolders (e.g., "AnnualCrop", "Forest", etc.) each containing images.

```

Listing 3: Custom Collate Function for DataLoader

```

1 from torch.utils.data import DataLoader
2 import torch
3
4 def collate_fn(batch):
5     # 'batch' is a list of tuples: (PIL_image, label_index).
6     images, labels = zip(*batch)
7     # Convert 'labels' (ints) to a single Tensor.
8     return list(images), torch.tensor(labels)

```

Listing 4: Dataset Splitting and Dataloader Setup

```

1 import torch

```

```

2 import torch.nn as nn
3 import torch.optim as optim
4 from torch.utils.data import DataLoader, SubsetRandomSampler
5 from torchvision import datasets
6 import numpy as np
7 import os
8 import matplotlib.pyplot as plt
9 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay,
    classification_report
10
11 from transformers import (
12     AutoImageProcessor,
13     AutoModelForImageClassification,
14     AutoModelForPreTraining,
15     AutoModel,
16     ViTMatteForImageMatting
17 )
18
19 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
20
21 data_dir = "/content/EuroSAT/2750"
22
23 # Create ImageFolder dataset from the unzipped directory
24 dataset = datasets.ImageFolder(root=data_dir)
25 classes = dataset.classes
26 num_classes = len(classes)
27 print("Classes:", classes)
28
29 # We'll define a validation split ratio of 0.2 (i.e., 20% for validation)
30 valid_size = 0.2
31 batch_size = 32
32
33 num_data = len(dataset)
34 indices = list(range(num_data))
35 np.random.shuffle(indices) # shuffle indices for random splitting
36 split = int(np.floor(valid_size * num_data))
37
38 # Partition dataset indices
39 train_idx, valid_idx = indices[split:], indices[:split]
40
41 train_sampler = SubsetRandomSampler(train_idx)
42 valid_sampler = SubsetRandomSampler(valid_idx)
43
44 # We define train_loader and valid_loader with collate_fn
45 train_loader = DataLoader(
46     dataset,
47     batch_size=batch_size,
48     sampler=train_sampler,
49     collate_fn=collate_fn
50 )
51
52 valid_loader = DataLoader(
53     dataset,
54     batch_size=batch_size,

```

```

55     sampler=valid_sampler,
56     collate_fn=collate_fn
57 )

```

Listing 5: Training, Validation, and Evaluation Routines

```

1  def train_one_epoch(model, processor, dataloader, optimizer):
2      """
3      Train the model for one epoch.
4
5      Args:
6          model: HF model for image classification
7          processor: HF image processor to transform PIL images to tensors
8          dataloader: DataLoader yielding (list_of_PILs, label_tensor)
9          optimizer: e.g. AdamW
10     """
11     model.train()
12     running_loss = 0.0
13     correct = 0
14     total = 0
15
16     for images, labels in dataloader:
17         labels = labels.to(device)
18         inputs = processor(images, return_tensors="pt").to(device)
19
20         outputs = model(**inputs, labels=labels)
21         loss = outputs["loss"]
22         logits = outputs["logits"]
23
24         optimizer.zero_grad()
25         loss.backward()
26         optimizer.step()
27
28         running_loss += loss.item() * labels.size(0)
29
30         _, predicted = logits.max(1)
31         total += labels.size(0)
32         correct += predicted.eq(labels).sum().item()
33
34     epoch_loss = running_loss / total
35     epoch_acc = 100.0 * correct / total
36     return epoch_loss, epoch_acc
37
38 def validate(model, processor, dataloader):
39     """
40     Validate model performance on a held-out set.
41     """
42     model.eval()
43     val_loss = 0.0
44     correct = 0
45     total = 0
46
47     with torch.no_grad():
48         for images, labels in dataloader:

```

```

49         labels = labels.to(device)
50         inputs = processor(images, return_tensors="pt").to(device)
51
52         outputs = model(**inputs, labels=labels)
53         loss = outputs["loss"]
54         logits = outputs["logits"]
55
56         val_loss += loss.item() * labels.size(0)
57
58         _, predicted = logits.max(1)
59         total += labels.size(0)
60         correct += predicted.eq(labels).sum().item()
61
62     val_loss /= total
63     val_acc = 100.0 * correct / total
64     return val_loss, val_acc
65
66 def evaluate_model(model, processor, dataloader):
67     """
68     Obtain all predictions and ground truths to compute classification metrics.
69     """
70     model.eval()
71     all_preds = []
72     all_labels = []
73
74     with torch.no_grad():
75         for images, labels in dataloader:
76             labels = labels.to(device)
77             inputs = processor(images, return_tensors="pt").to(device)
78             outputs = model(**inputs)
79             logits = outputs["logits"]
80
81             _, predicted = logits.max(1)
82             all_preds.extend(predicted.cpu().numpy())
83             all_labels.extend(labels.cpu().numpy())
84
85     return all_labels, all_preds
86
87 def plot_confusion_matrix(all_labels, all_preds, classes):
88     """
89     Plot a confusion matrix using sklearn's ConfusionMatrixDisplay.
90     """
91     cm = confusion_matrix(all_labels, all_preds)
92     disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=classes)
93     disp.plot(cmap=plt.cm.Blues, xticks_rotation='vertical', values_format='d')
94     plt.title("Confusion Matrix on Validation Set")
95     plt.show()
96
97 def print_classification_report(all_labels, all_preds, classes):
98     """
99     Print a detailed classification report with precision, recall, F1.
100     """
101     report = classification_report(all_labels, all_preds, target_names=classes)
102     print("Classification Report:\n", report)

```

Listing 6: High-Level Training and Evaluation Loop

```

1 def train_and_evaluate(model, processor, train_loader, valid_loader, epochs=35, lr=1e-3,
2   wd=1e-4):
3     """
4     Train the model for several epochs, track best val accuracy, and then
5     produce final plots and classification metrics.
6
7     Args:
8         model: HF model for classification
9         processor: HF image processor
10        train_loader, valid_loader: DataLoaders
11        epochs: number of epochs
12        lr: learning rate
13        wd: weight decay
14    """
15    optimizer = optim.AdamW(model.parameters(), lr=lr, weight_decay=wd)
16    scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=10)
17
18    best_val_acc = 0.0
19    train_losses, train_accuracies = [], []
20    val_losses, val_accuracies = [], []
21
22    best_model_state = None
23
24    for epoch in range(epochs):
25        train_loss, train_acc = train_one_epoch(model, processor, train_loader, optimizer)
26        val_loss, val_acc = validate(model, processor, valid_loader)
27
28        scheduler.step()
29
30        train_losses.append(train_loss)
31        train_accuracies.append(train_acc)
32        val_losses.append(val_loss)
33        val_accuracies.append(val_acc)
34
35        print(f"Epoch [{epoch+1}/{epochs}]")
36        print(f"Train Loss: {train_loss:.4f}, Train Acc: {train_acc:.2f}%")
37        print(f"Val Loss: {val_loss:.4f}, Val Acc: {val_acc:.2f}%\n")
38
39        # Track the best model
40        if val_acc > best_val_acc:
41            best_val_acc = val_acc
42            best_model_state = model.state_dict()
43
44    if best_model_state is not None:
45        model.load_state_dict(best_model_state)
46
47    # Plot training vs. validation loss
48    plt.figure(figsize=(10,5))
49    plt.title("Training and Validation Loss")
50    plt.plot(train_losses, label="Train Loss")
51    plt.plot(val_losses, label="Validation Loss")

```

```

51 plt.xlabel("Epochs")
52 plt.ylabel("Loss")
53 plt.legend()
54 plt.show()
55
56 # Plot training vs. validation accuracy
57 plt.figure(figsize=(10,5))
58 plt.title("Training and Validation Accuracy")
59 plt.plot(train_accuracies, label="Train Accuracy")
60 plt.plot(val_accuracies, label="Validation Accuracy")
61 plt.xlabel("Epochs")
62 plt.ylabel("Accuracy (%)")
63 plt.legend()
64 plt.show()
65
66 # Final evaluation on validation set
67 all_labels, all_preds = evaluate_model(model, processor, valid_loader)
68 plot_confusion_matrix(all_labels, all_preds, classes)
69 print_classification_report(all_labels, all_preds, classes)
70
71 return best_val_acc

```

Listing 7: Loading and Training the DeiT Model

```

1 def get_deit_model(num_classes):
2     """
3     Load the 'facebook/deit-small-patch16-224' model & processor for classification.
4     """
5     from transformers import AutoImageProcessor, AutoModelForImageClassification
6     processor = AutoImageProcessor.from_pretrained(
7         "facebook/deit-small-patch16-224",
8         use_fast=True
9     )
10    model = AutoModelForImageClassification.from_pretrained(
11        "facebook/deit-small-patch16-224",
12        num_labels=num_classes,
13        ignore_mismatched_sizes=True
14    )
15    model.to(device)
16    return model, processor
17
18 print("=== DeiT ===")
19 deit_model, deit_processor = get_deit_model(num_classes)
20 deit_acc = train_and_evaluate(deit_model, deit_processor, train_loader, valid_loader)
21 print(f"DeiT Val Acc: {deit_acc:.2f}%")

```

Listing 8: Loading and Training the Swin Transformer Model

```

1 def get_swin_model(num_classes):
2     """
3     Load the 'microsoft/swin-tiny-patch4-window7-224' Swin Transformer for classification
4     """
5     from transformers import AutoImageProcessor, AutoModelForImageClassification

```

```

6     processor = AutoImageProcessor.from_pretrained(
7         "microsoft/swin-tiny-patch4-window7-224",
8         use_fast=True
9     )
10    model = AutoModelForImageClassification.from_pretrained(
11        "microsoft/swin-tiny-patch4-window7-224",
12        num_labels=num_classes,
13        ignore_mismatched_sizes=True
14    )
15    model.to(device)
16    return model, processor
17
18    print("=== Swin Transformer ===")
19    swin_model, swin_processor = get_swin_model(num_classes)
20    swin_acc = train_and_evaluate(swin_model, swin_processor, train_loader, valid_loader)
21    print(f"Swin Transformer Val Acc: {swin_acc:.2f}%")

```

Listing 9: Loading and Training the MAE-based Model

```

1  def get_mae_model(num_classes):
2      """
3      Load a pretrained MAE (Masked Autoencoder) model from 'facebook/vit-mae-large'
4      and adapt it for classification by adding a custom head.
5      """
6      from transformers import AutoImageProcessor, AutoModelForPreTraining
7      processor = AutoImageProcessor.from_pretrained("facebook/vit-mae-large", use_fast=
8          True)
9      pretrained_model = AutoModelForPreTraining.from_pretrained("facebook/vit-mae-large")
10     base_model = pretrained_model.vit
11
12     classification_head = nn.Linear(base_model.config.hidden_size, num_classes)
13
14     class MAEForClassification(nn.Module):
15         def __init__(self, base_model, classification_head):
16             super().__init__()
17             self.base_model = base_model
18             self.classifier = classification_head
19
20         def forward(self, pixel_values, labels=None):
21             outputs = self.base_model(pixel_values=pixel_values)
22             pooled_output = outputs.last_hidden_state[:, 0, :]
23             logits = self.classifier(pooled_output)
24             loss = None
25             if labels is not None:
26                 loss_fn = nn.CrossEntropyLoss()
27                 loss = loss_fn(logits, labels)
28             return {"loss": loss, "logits": logits}
29
30     model = MAEForClassification(base_model, classification_head).to(device)
31     return model, processor
32
33    print("=== MAE ===")
34    mae_model, mae_processor = get_mae_model(num_classes)
35    mae_acc = train_and_evaluate(mae_model, mae_processor, train_loader, valid_loader)

```



```
35 print(f"MAE Val Acc: {mae_acc:.2f}%")
```

Listing 10: Loading and Training the PVT (Pyramid Vision Transformer) Model

```
1 def get_pvt_model(num_classes):
2     """
3     Load 'Zetatech/pvt-tiny-224' for classification, specifying 'num_labels=num_classes'.
4     """
5     import torch
6     import torch.nn as nn
7     from transformers import AutoImageProcessor, PvtForImageClassification
8
9     device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
10    processor = AutoImageProcessor.from_pretrained("Zetatech/pvt-tiny-224")
11    model = PvtForImageClassification.from_pretrained(
12        "Zetatech/pvt-tiny-224",
13        num_labels=num_classes,
14        problem_type="single_label_classification",
15        ignore_mismatched_sizes=True
16    )
17    model.to(device)
18    return model, processor
19
20 print("=== PVT ===")
21 pvt_model, pvt_processor = get_pvt_model(num_classes)
22 pvt_acc = train_and_evaluate(pvt_model, pvt_processor, train_loader, valid_loader)
23 print(f"PVT Validation Accuracy: {pvt_acc:.2f}%")
```

Listing 11: Loading and Training the YOLOs Model

```
1 def get_yolos_small_for_classification(num_classes):
2     """
3     Convert the 'hustvl/yolos-small' model (normally for object detection)
4     into a classification model by adding a linear head.
5     """
6     import torch
7     import torch.nn as nn
8     from transformers import YolosModel, AutoImageProcessor
9
10    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
11    processor = AutoImageProcessor.from_pretrained("hustvl/yolos-small")
12    base_model = YolosModel.from_pretrained("hustvl/yolos-small")
13
14    hidden_size = base_model.config.hidden_size
15    classification_head = nn.Linear(hidden_size, num_classes)
16
17    class YOLOSSmallForClassification(nn.Module):
18        def __init__(self, base_model, classifier):
19            super().__init__()
20            self.base_model = base_model
21            self.classifier = classifier
22
23        def forward(self, pixel_values, labels=None):
24            outputs = self.base_model(pixel_values=pixel_values, return_dict=True)
```

```

25         pooled_output = outputs.pooler_output # YOLOs model's global feature
26         logits = self.classifier(pooled_output)
27         loss = None
28         if labels is not None:
29             loss_fn = nn.CrossEntropyLoss()
30             loss = loss_fn(logits, labels)
31         return {"loss": loss, "logits": logits}
32
33     model = YOLOSmallForClassification(base_model, classification_head).to(device)
34     return model, processor
35
36 print("=== YOLOs ===")
37 yolos_model, yolos_processor = get_yolos_small_for_classification(num_classes=num_classes
38 )
39 yolos_acc = train_and_evaluate(yolos_model, yolos_processor, train_loader, valid_loader)
40 print(f"YOLOs Validation Accuracy: {yolos_acc:.2f}%")

```